

# imitation learning

## behavior cloning (lecture2 part2)

定义：直接使用监督学习的方法，将专家数据看作样本输入

### 目标

我们在 $p_{\text{data}}(\mathbf{o}_t)$ 下训练，但是我们用 $p_{\pi_\theta}(\mathbf{o}_t)$ 测试，但 $p_{\text{data}}(\mathbf{o}_t) \neq p_{\pi_\theta}(\mathbf{o}_t)$

所以不考虑以下目标函数

$$\max_{\theta} E_{\mathbf{o}_t \sim p_{\text{data}}(\mathbf{o}_t)} [\log \pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)]$$

首先，定义损失函数为

$$c(\mathbf{s}_t, \mathbf{a}_t) = \begin{cases} 0 & \text{if } \mathbf{a}_t = \pi^*(\mathbf{s}_t) \\ 1 & \text{otherwise} \end{cases}$$

我们关心的是 $\pi_\theta$ 情况，我们需要做的是

$$\text{minimize } E_{\mathbf{s}_t \sim p_{\pi_\theta}(\mathbf{s}_t)} [c(\mathbf{s}_t, \mathbf{a}_t)]$$

假设： $\pi_\theta(\mathbf{a} \neq \pi^*(\mathbf{s}) | \mathbf{s}) \leq \epsilon$ ，对于 $\mathbf{s} \in \mathcal{D}_{\text{train}}$

因此有

$$E \left[ \sum_t c(\mathbf{s}_t, \mathbf{a}_t) \right] \leq \underbrace{\epsilon T + (1 - \epsilon)(\epsilon(T - 1) + (1 - \epsilon)(\dots))}_{T \text{ terms, each } O(\epsilon T)} = O(\epsilon T^2)$$

这个式子表明BC的cost会随着决策步数的增加而呈现平方次增加

更加泛化的说，对于 $\mathbf{s} \sim p_{\text{train}}(\mathbf{s})$ ，我们假设 $E_{p_{\text{train}}(\mathbf{s})} [\pi_\theta(\mathbf{a} \neq \pi^*(\mathbf{s}) | \mathbf{s})] \leq \epsilon$

对于 $p_{\text{train}}(\mathbf{s}) \neq p_\theta(\mathbf{s})$ ，我们有

$$p_\theta(\mathbf{s}_t) = (1 - \epsilon)^t p_{\text{train}}(\mathbf{s}_t) + (1 - (1 - \epsilon)^t) p_{\text{mistake}}(\mathbf{s}_t)$$

式子前一项表明不犯错（即train data）的概率，后一项表明其他分布

然后有

$$|p_\theta(\mathbf{s}_t) - p_{\text{train}}(\mathbf{s}_t)| = (1 - (1 - \epsilon)^t) |p_{\text{mistake}}(\mathbf{s}_t) - p_{\text{train}}(\mathbf{s}_t)| \leq 2(1 - (1 - \epsilon)^t)$$

其中 $(1 - \epsilon)^t \geq 1 - \epsilon t$ ，for  $\epsilon \in [0, 1]$ ，所以小于 $2\epsilon t$

对于我们的目标函数来说

$$\begin{aligned} \sum_t E_{p_\theta(\mathbf{s}_t)} [c_t] &= \sum_t \sum_{\mathbf{s}_t} p_\theta(\mathbf{s}_t) c_t(\mathbf{s}_t) \leq \sum_t \sum_{\mathbf{s}_t} p_{\text{train}}(\mathbf{s}_t) c_t(\mathbf{s}_t) + |p_\theta(\mathbf{s}_t) - p_{\text{train}}(\mathbf{s}_t)| c_{\text{max}} \\ &\leq \sum_t \epsilon + 2\epsilon t \leq \epsilon T + 2\epsilon T^2 \\ &O(\epsilon T^2) \end{aligned}$$

# problem of Imitation learning

## Non-Markovian behavior

人类在标注数据进行决策的时候，考虑到的并不仅仅是当前这一帧的observation，而是根据过去的许多信息来做出判断的，对于同样一帧图像，可能由于前面的信息不同，人类会做出不一样的决策。对于这样的数据，算法在学习的时候是有可能产生混淆的，这也就导致算法对于数据的拟合效果不够好。

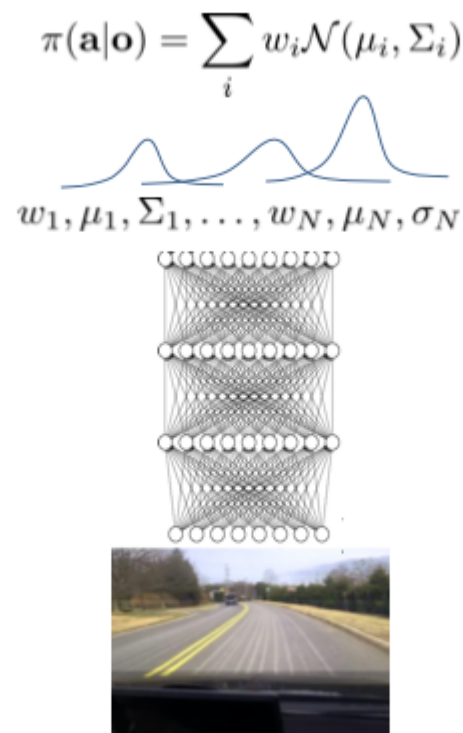
所以解决问题的关键就是如何将缺失的信息补充回来。最直观的方法是不仅仅使用当前帧作为输入，而是使用过去若干帧作为输入并从中提取有用的信息，这也正是原始DQN中使用的方法。RNN是一个不得不提到的方法，它将前面帧的信息作为hidden state传输给后面的状态，从而作为缺失信息的补充。

但是**添加的信息越多则更可能导致泛化性能变差**。但是学习中由于信息过多，加上现有的学习方法大部分对于因果关系的提取仍然比较差，所以会将这种共现因子当作是决策的原因，从而产生了所谓的causal confusion


## Multimodal behavior

在很多时候对于一个情景的确是存在多个可行解


## Mixture of Gaussian



Latent variable model



Autoregressive discretization



Extra notes

连续动作vs离散动作空间

特征	连续动作空间 ( Box )	离散动作空间 ( Discrete )
动作类型	实数向量	整数编号
Gym 类型	<code>gym.spaces.Box</code>	<code>gym.spaces.Discrete</code>

特征	连续动作空间 ( Box )	离散动作空间 ( Discrete )
策略分布类型	<code>Normal(mean, std)</code>	<code>Categorical(probs)</code>
适用算法举例	DDPG, PPO (continuous), SAC	DQN, REINFORCE, PPO (discrete)
处理方式 (log_prob)	<code>dist.log_prob(actions).sum(-1)</code>	<code>dist.log_prob(action)</code>

# homework

## analysis

### 第一问

对于其中任意某个状态来说，有

$$\mathbb{E}_{p_{\pi^*}(s)} \pi_{\theta}(a \neq \pi^*(s) \mid s) = \frac{1}{T} \sum_{t=1}^T \mathbb{E}_{p_{\pi^*}(s_t)} \pi_{\theta}(a_t \neq \pi^*(s_t) \mid s_t) \leq \varepsilon$$

这个式子表明在状态 $s_t$ 下policy $\pi_{\theta}$  dispute policy  $\pi^*$ 的概率小于 $\varepsilon$

用题目提示的公式有

$$\Pr \left( \bigcup_{t=1}^T E_t \right) \leq \sum_{t=1}^T \Pr(E_t) \leq T\epsilon$$

其中 $E_t$ 代表the event that  $\pi_{\theta}$  dispute with  $\pi^*$  at time t

故得(a conservative bound)

$$\sum_{s_t} |p_{\pi_{\theta}}(s_t) - p_{\pi^*}(s_t)| \leq 2T\varepsilon$$

### 第二问

#### 2.a

因为reward only depend on the last state,故有

$$J(\pi^*) - J(\pi_{\theta}) \leq R_{\max} T \epsilon = O(T\epsilon)$$

#### 2.b

根据题目所给公式 $J(\pi) = \sum_{t=1}^T \mathbb{E}_{p_{\pi}(s_t)} r(s_t)$ ,

代入可得

$$J(\pi^*) - J(\pi_{\theta}) \leq R_{\max} \sum_{t=1}^T t\epsilon = R_{\max} \frac{T(T+1)}{2} \epsilon = O(T^2\epsilon)$$

## Editing code

- 在安装环境时，如果直接 `pip install -r requirements.txt`，会出现报错无法安装box2d-py依赖项，可以先通过 `conda install conda-forge::box2d-py` 安装box2d-py,再 `pip install -r requirements.txt`

- ```
1 | INTEL oneMKL ERROR: 找不到指定的模块。 mkl_intel_thread.2.dll.
2 | Intel oneMKL FATAL ERROR: Cannot load mkl_intel_thread.2.dll.
```

后续报错找不到numpy，查找到应该mkl，mkl-service没有正确安装

```
conda install mkl
```

```
pip install mkl-service
```

至此环境搭配完毕

## MLP\_policy.py

### forward

- 给定 observation（状态），**输出一个分布**，表示该状态下采取各个动作的概率。

```
1 | def forward(self, observation: torch.FloatTensor) -> Any:
2 |     """
3 |     Defines the forward pass of the network
4 |
5 |     :param observation: observation(s) to query the policy
6 |     :return:
7 |         action: sampled action(s) from the policy
8 |     """
9 |     # TODO: implement the forward pass of the network.
10 |    # You can return anything you want, but you should be able to differentiate
11 |    # through it. For example, you can return a torch.FloatTensor. You can also
12 |    # return more flexible objects, such as a
13 |    # `torch.distributions.Distribution` object. It's up to you!
14 |    observation=observation.to(ptu.device)
15 |    mean = self.mean_net(observation)#调用mean_net进行前向传播
16 |    logstd = self.logstd.expand_as(mean) # logstd is a single value
17 |    std = torch.exp(logstd) # convert logstd to std
18 |    dist = distributions.Normal(mean, std) # create a normal distribution
19 |    action = dist.rsample() # sample an action from the distribution,使用rsample可
    以使得梯度可以通过采样传递
20 |    action = torch.tanh(action) # apply tanh to the action限制
21 |    return action, dist
```

- observation是一个 **PyTorch 张量**（`torch.FloatTensor`），形状是（batch\_size,ob\_dim）
- mean\_net是一个神经网络，输出是动作分布的均值
- [torch.Tensor.expand\\_as — PyTorch 2.7 documentation](#)：将logstd拓展成mean一样的大小
- [Probability distributions - torch.distributions — PyTorch 2.7 documentation](#)：.Normal创建一个多维独立正态分布

$$\pi(a | s) = \mathcal{N}(\mu(s), \sigma(s)^2)$$

- 使用 `rsample()` 而不是 `sample()` 是为了支持 **重参数化 (reparameterization trick)**，即允许梯度从 loss 反传到 `mean` 和 `logstd`，常用于策略梯度或 VAE。
- 将动作经过 `tanh` 激活函数映射到 `(-1, 1)` 范围，使得输出动作符合大多数连续控制环境（如 MuJoCo）要求的动作空间限制。

## update

- 给定一批 `(observation, action)` 样本（专家数据），
- 调用 `forward` 得到策略在这些状态下的动作分布；
- 计算这些 **分布对 expert 动作的 log-probability**；
- 最大化该 log-prob（或最小化负对数似然）作为 loss，反向传播并优化网络。

```

1  def update(self, observations, actions):
2      """
3      Updates/trains the policy
4
5      :param observations: observation(s) to query the policy
6      :param actions: actions we want the policy to imitate
7      :return:
8          dict: 'Training Loss': supervised learning loss
9      """
10     # TODO: update the policy and return the loss
11     observations = observations.to(ptu.device)
12     actions = actions.to(ptu.device)
13     self.optimizer.zero_grad()
14     _, dist = self.forward(observations)
15     # Calculate the loss as the negative log likelihood of the actions under the
policy
16     log_probs = dist.log_prob(actions)
17     log_probs = log_probs.sum(dim=-1) # sum over action dimensions
18     loss = -log_probs.mean() # mean over batch
19     # Backpropagate the loss
20     loss.backward()
21     self.optimizer.step()
22     # Return a dictionary with the training loss
23     # You can add extra logging information here, but keep this line
24     return {
25         # You can add extra logging information here, but keep this line
26         'Training Loss': ptu.to_numpy(loss),
27     }

```

- `self.optimizer.zero_grad()` 梯度清零，为当前batch的梯度计算做准备
- 计算专家动作 `actions` 在当前策略分布下的对数概率（log likelihood），如果动作是多维的，求和得到整体动作向量的对数概率 `.sum(dim=-1)`
- `loss = -log_probs.mean()` 计算负对数似然损失（negative log likelihood, NLL）  
损失越小，说明策略分布越能“解释”专家动作，即策略输出的动作更接近专家动作
- `loss.backward()` 反向计算梯度
- `self.optimizer.step()` 优化器根据梯度更新模型参数，提升策略

## utils.py

### sample\_trajectory

- 用当前策略采取动作，并记录每个时间步的observation、action、reward、next\_observation、terminal

```
1 def sample_trajectory(env, policy, max_path_length, render=False):
2     """Sample a rollout in the environment from a policy."""
3
4     # initialize env for the beginning of a new rollout
5     ob = env.reset() # TODO: initial observation after resetting the env
6     print("observation shape: ", ob.shape)
7     # init vars
8     obs, acs, rewards, next_obs, terminals, image_obs = [], [], [], [], [], []
9     steps = 0
10    while True:
11
12        # render image of the simulated env
13        if render:
14            if hasattr(env, 'sim'):
15                img = env.sim.render(camera_name='track', height=500, width=500)[::-1]
16            else:
17                img = env.render(mode='single_rgb_array')
18            image_obs.append(cv2.resize(img, dsize=(250, 250),
19 interpolation=cv2.INTER_CUBIC))
20
21        # TODO use the most recent ob to decide what to do
22        ac, _ = policy.forward(ptu.from_numpy(ob)) # HINT: this is a numpy array
23        ac = ac[0] #取出单个动作, 去除batch维度
24
25        # TODO: take that action and get reward and next ob
26        next_ob, rew, done, _ = env.step(ac) # HINT: this is a numpy array
27
28        # TODO rollout can end due to done, or due to max_path_length
29        steps += 1
30        rollout_done = done or steps >= max_path_length # HINT: this is either 0 or 1
31
32        # record result of taking that action
33        obs.append(ob)
34        acs.append(ac)
35        rewards.append(rew)
36        next_obs.append(next_ob)
37        terminals.append(rollout_done)
38
39        ob = next_ob # jump to next timestep
40
41        # end the rollout if the rollout ended
42        if rollout_done:
43            break
44
45    return {"observation" : np.array(obs, dtype=np.float32),
46            "image_obs" : np.array(image_obs, dtype=np.uint8),
47            "reward" : np.array(rewards, dtype=np.float32),
```

```

47         "action" : np.array(acs, dtype=np.float32),
48         "next_observation": np.array(next_obs, dtype=np.float32),
49         "terminal": np.array(terminals, dtype=np.float32)}

```

- `if hasattr(env, 'sim')` 检查环境是否有sim属性，有的话通过`sim.render`渲染，`[::-1]`将图像上下翻转，因为MuJoCo 渲染默认是上下颠倒的，`cv2.INTER_CUBIC` (Bicubic Interpolation) 双三次插值
- `env.step(ac)` 让环境执行一个动作，并返回它对这个动作的“回应”

### sample\_trajectories和sample\_n\_trajectories

```

1  def sample_trajectories(env, policy, min_timesteps_per_batch, max_path_length,
2      render=False):
3      """Collect rollouts until we have collected min_timesteps_per_batch steps."""
4
5      timesteps_this_batch = 0
6      paths = []
7      while timesteps_this_batch < min_timesteps_per_batch:
8
9          #collect rollout
10         path = sample_trajectory(env, policy, max_path_length, render)
11         paths.append(path)
12
13         #count steps
14         timesteps_this_batch += get_pathlength(path)
15
16     return paths, timesteps_this_batch

```

- 多次调用 `sample_trajectory`，直到采集的时间步 (step) 数 ≥ `min_timesteps_per_batch`, `max_path_length` 是单条路径的最大步数，`paths` 是所有采集到的 trajectories 列表

同理有

```

1  def sample_n_trajectories(env, policy, ntraj, max_path_length, render=False):
2      """Collect ntraj rollouts."""
3
4      paths = []
5      for i in range(ntraj):
6          # collect rollout
7          path = sample_trajectory(env, policy, max_path_length, render)
8          paths.append(path)
9      return paths

```

### convert\_listofrollouts

将这些多个轨迹中的对应数据拼接 (concatenate) 起来，变成一个大数组，方便后续处理

```

1  def convert_listofrollouts(paths, concat_rew=True):
2      """
3          Take a list of rollout dictionaries
4          and return separate arrays,

```



```

5         where each array is a concatenation of that array from across the rollouts
6         """
7         observations = np.concatenate([path["observation"] for path in paths])
8         actions = np.concatenate([path["action"] for path in paths])
9         if concat_rew:
10             rewards = np.concatenate([path["reward"] for path in paths])
11         else:
12             rewards = [path["reward"] for path in paths]
13         next_observations = np.concatenate([path["next_observation"] for path in paths])
14         terminals = np.concatenate([path["terminal"] for path in paths])
15         return observations, actions, rewards, next_observations, terminals

```

- `concat_rew=True` — 返回所有轨迹奖励的拼接数组（扁平的，合并的） `concat_rew=False` — 返回奖励的列表（每个元素是单个轨迹的奖励序列）

## compute\_metrics

根据训练轨迹和评估轨迹（即 `paths` 和 `eval_paths`）计算训练日志中常见的一些评估指标

```

1 def compute_metrics(paths, eval_paths):
2     """Compute metrics for logging."""
3
4     # returns, for logging
5     train_returns = [path["reward"].sum() for path in paths]
6     eval_returns = [eval_path["reward"].sum() for eval_path in eval_paths]
7
8     # episode lengths, for logging
9     train_ep_lens = [len(path["reward"]) for path in paths]
10    eval_ep_lens = [len(eval_path["reward"]) for eval_path in eval_paths]
11
12    # decide what to log
13    logs = OrderedDict()
14    logs["Eval_AverageReturn"] = np.mean(eval_returns)
15    logs["Eval_StdReturn"] = np.std(eval_returns)
16    logs["Eval_MaxReturn"] = np.max(eval_returns)
17    logs["Eval_MinReturn"] = np.min(eval_returns)
18    logs["Eval_AverageEpLen"] = np.mean(eval_ep_lens)
19
20    logs["Train_AverageReturn"] = np.mean(train_returns)
21    logs["Train_StdReturn"] = np.std(train_returns)
22    logs["Train_MaxReturn"] = np.max(train_returns)
23    logs["Train_MinReturn"] = np.min(train_returns)
24    logs["Train_AverageEpLen"] = np.mean(train_ep_lens)
25
26    return logs

```

- 使用 `OrderedDict` 是为了让输出的指标顺序稳定（可选，不影响结果，但便于日志文件查看）

## run\_hw1

行为克隆 (Behavior Cloning) 与 DAgger 的主训练脚本，作用是训练一个模仿学习的策略网络(核心)

tensorboard --logdir=你的日志目录路径，查看具体日志

```
1  # how many rollouts to save as videos to tensorboard
2  MAX_NVIDEO = 2#每次迭代最多保存两个视频
3  MAX_VIDEO_LEN = 40 # we overwrite this in the code below
4
5  MJ_ENV_NAMES = ["Ant-v4", "walker2d-v4", "HalfCheetah-v4", "Hopper-v4"]
6
7
8  def run_training_loop(params):
9      """
10     Runs training with the specified parameters
11     (behavior cloning or dagger)
12
13     Args:
14         params: experiment parameters
15     """
16
17     #####
18     ## INIT
19     #####
20
21     # Get params, create logger, create TF session
22     logger = Logger(params['logdir'])
23
24     # Set random seeds
25     seed = params['seed']
26     np.random.seed(seed)
27     torch.manual_seed(seed)
28     ptu.init_gpu(
29         use_gpu=not params['no_gpu'],
30         gpu_id=params['which_gpu']
31     )
32
33     # Set logger attributes
34     log_video = True
35     log_metrics = True
36
37     #####
38     ## ENV
39     #####
40
41     # Make the gym environment
42     env = gym.make(params['env_name'], render_mode=None)
43     env.reset(seed=seed)
44
45     # Maximum length for episodes
46     params['ep_len'] = params['ep_len'] or env.spec.max_episode_steps
47     MAX_VIDEO_LEN = params['ep_len']
```

```

48
49     assert isinstance(env.action_space, gym.spaces.Box), "Environment must be
continuous"#连续动作空间
50     # Observation and action sizes
51     ob_dim = env.observation_space.shape[0]
52     ac_dim = env.action_space.shape[0]
53
54     # simulation timestep, will be used for video saving
55     if 'model' in dir(env):
56         fps = 1/env.model.opt.timestep
57     else:
58         fps = env.env.metadata['render_fps']
59
60     #####
61     ## AGENT
62     #####
63
64     # TODO: Implement missing functions in this class.
65     actor = MLPPolicySL(
66         ac_dim,
67         ob_dim,
68         params['n_layers'],
69         params['size'],
70         learning_rate=params['learning_rate'],
71     )
72
73     # replay buffer
74     replay_buffer = ReplayBuffer(params['max_replay_buffer_size'])
75
76     #####
77     ## LOAD EXPERT POLICY
78     #####
79
80     print('Loading expert policy from...', params['expert_policy_file'])
81     expert_policy = LoadedGaussianPolicy(params['expert_policy_file'])
82     expert_policy.to(ptu.device)
83     print('Done restoring expert policy...')
84
85     #####
86     ## TRAINING LOOP
87     #####
88
89     # init vars at beginning of training
90     total_envsteps = 0
91     start_time = time.time()
92
93     for itr in range(params['n_iter']):
94         print("\n\n***** Iteration %i *****"%itr)
95
96         # decide if videos should be rendered/logged at this iteration
97         log_video = ((itr % params['video_log_freq'] == 0) and
(params['video_log_freq'] != -1))
98         # decide if metrics should be logged

```

```

99     log_metrics = (itr % params['scalar_log_freq'] == 0)
100
101     print("\nCollecting data to be used for training...")
102     if itr == 0:
103         # BC training from expert data.
104         paths = pickle.load(open(params['expert_data'], 'rb'))
105         envsteps_this_batch = 0
106     else:
107         # DAGGER training from sampled data relabeled by expert
108         assert params['do_dagger']
109         # TODO: collect `params['batch_size']` transitions
110         # HINT: use utils.sample_trajectories
111         # TODO: implement missing parts of utils.sample_trajectory
112         paths, envsteps_this_batch = utils.sample_trajectories(
113             env, actor, params['batch_size'], params['ep_len'], render=log_video)
114
115         # relabel the collected obs with actions from a provided expert policy
116         if params['do_dagger']:
117             print("\nRelabelling collected observations with labels from an expert
118 policy...")
119
120             # TODO: relabel collected observations (from our policy) with labels
121             # HINT: query the policy (using the get_action function) with paths[i]
122             # and replace paths[i]["action"] with these expert labels
123             for i in range(len(paths)):
124                 paths[i]["action"] = expert_policy.get_action(
125                     paths[i]["observation"])
126                 #利用专家数据替代
127
128             total_envsteps += envsteps_this_batch
129             # add collected data to replay buffer
130             replay_buffer.add_rollouts(paths)
131
132             # train agent (using sampled data from replay buffer)
133             print('\nTraining agent using sampled data from replay buffer...')
134             training_logs = []
135             for _ in range(params['num_agent_train_steps_per_iter']):
136
137                 # TODO: sample some data from replay_buffer
138                 # HINT1: how much data = params['train_batch_size']
139                 # HINT2: use np.random.permutation to sample random indices
140                 # HINT3: return corresponding data points from each array (i.e., not
141                 # different indices from each array)
142                 # for imitation learning, we only need observations and actions.
143                 ob_batch, ac_batch = replay_buffer.sample(
144                     params['train_batch_size'])
145
146                 # use the sampled data to train an agent
147                 train_log = actor.update(ob_batch, ac_batch)
148                 training_logs.append(train_log)

```

```

148     # log/save
149     print('\nBeginning logging procedure...')
150     if log_video:
151         # save eval rollouts as videos in tensorboard event file
152         print('\nCollecting video rollouts eval')
153         eval_video_paths = utils.sample_n_trajectories(
154             env, actor, MAX_NVIDEO, MAX_VIDEO_LEN, True)
155
156         # save videos
157         if eval_video_paths is not None:
158             logger.log_paths_as_videos(
159                 eval_video_paths, itr,
160                 fps=fps,
161                 max_videos_to_save=MAX_NVIDEO,
162                 video_title='eval_rollouts')
163
164     if log_metrics:
165         # save eval metrics
166         print("\nCollecting data for eval...")
167         eval_paths, eval_envsteps_this_batch = utils.sample_trajectories(
168             env, actor, params['eval_batch_size'], params['ep_len'])
169
170         logs = utils.compute_metrics(paths, eval_paths)
171         # compute additional metrics
172         logs.update(training_logs[-1]) # Only use the last log for now
173         logs["Train_EnvstepsSoFar"] = total_envsteps
174         logs["TimeSinceStart"] = time.time() - start_time
175         if itr == 0:
176             logs["Initial_DataCollection_AverageReturn"] =
logs["Train_AverageReturn"]
177
178         # perform the logging
179         for key, value in logs.items():
180             print('{} : {}'.format(key, value))
181             logger.log_scalar(value, key, itr)
182         print('Done logging...\n\n')
183
184         logger.flush()
185
186     if params['save_params']:
187         print('\nSaving agent params')
188         actor.save('{} /policy_itr_{}.pt'.format(params['logdir'], itr))
189
190
191 def main():
192     import argparse
193     parser = argparse.ArgumentParser()
194     parser.add_argument('--expert_policy_file', '-epf', type=str, required=True) #
relative to where you're running this script from
195     parser.add_argument('--expert_data', '-ed', type=str, required=True) #relative to
where you're running this script from
196     parser.add_argument('--env_name', '-env', type=str, help=f'choices: {"',
".join(MJ_ENV_NAMES)}", required=True)

```

```

197     parser.add_argument('--exp_name', '-exp', type=str, default='pick an experiment
name', required=True)
198     parser.add_argument('--do_dagger', action='store_true')
199     parser.add_argument('--ep_len', type=int)
200
201     parser.add_argument('--num_agent_train_steps_per_iter', type=int, default=1000) #
number of gradient steps for training policy (per iter in n_iter)
202     parser.add_argument('--n_iter', '-n', type=int, default=1)
203
204     parser.add_argument('--batch_size', type=int, default=1000) # training data
collected (in the env) during each iteration
205     parser.add_argument('--eval_batch_size', type=int,
206                         default=1000) # eval data collected (in the env) for logging
metrics
207     parser.add_argument('--train_batch_size', type=int,
208                         default=100) # number of sampled data points to be used per
gradient/train step
209
210     parser.add_argument('--n_layers', type=int, default=2) # depth, of policy to be
learned
211     parser.add_argument('--size', type=int, default=64) # width of each layer, of
policy to be learned
212     parser.add_argument('--learning_rate', '-lr', type=float, default=5e-3) # LR for
supervised learning
213
214     parser.add_argument('--video_log_freq', type=int, default=5)
215     parser.add_argument('--scalar_log_freq', type=int, default=1)
216     parser.add_argument('--no_gpu', '-ngpu', action='store_true')
217     parser.add_argument('--which_gpu', type=int, default=0)
218     parser.add_argument('--max_replay_buffer_size', type=int, default=1000000)
219     parser.add_argument('--save_params', action='store_true')
220     parser.add_argument('--seed', type=int, default=1)
221     args = parser.parse_args()
222
223     # convert args to dictionary
224     params = vars(args)
225
226     #####
227     ### CREATE DIRECTORY FOR LOGGING
228     #####
229
230     if args.do_dagger:
231         # Use this prefix when submitting. The auto-grader uses this prefix.
232         logdir_prefix = 'q2_'
233         assert args.n_iter>1, ('DAGGER needs more than 1 iteration (n_iter>1) of
training, to iteratively query the expert and train (after 1st warmstarting from
behavior cloning).')
234     else:
235         # Use this prefix when submitting. The auto-grader uses this prefix.
236         logdir_prefix = 'q1_'
237         assert args.n_iter==1, ('Vanilla behavior cloning collects expert data just
once (n_iter=1)')
238

```

```

239     # directory for logging
240     data_path = os.path.join(os.path.dirname(os.path.realpath(__file__)),
241                               '../..data')
241     if not (os.path.exists(data_path)):
242         os.makedirs(data_path)
243     logdir = logdir_prefix + args.exp_name + '_' + args.env_name + '_' +
244             time.strftime("%d-%m-%Y_%H-%M-%S")
245     logdir = os.path.join(data_path, logdir)
246     params['logdir'] = logdir
247     if not(os.path.exists(logdir)):
248         os.makedirs(logdir)
249
250     #####
251     ### RUN TRAINING
252     #####
253     run_training_loop(params)
254
255 if __name__ == "__main__":
256     main()
257
258

```

- `logger = Logger(params['logdir'])` 通过Logger类设置日志目录

```

1 class Logger:
2     def __init__(self, log_dir, n_logged_samples=10, summary_writer=None):
3         self._log_dir = log_dir
4         print('#####')
5         print('logging outputs to ', log_dir)
6         print('#####')
7         self._n_logged_samples = n_logged_samples
8         self._summ_writer = Summarywriter(log_dir, flush_secs=1, max_queue=1)
9

```

这个是 `tensorboardx.Summarywriter` 的实例，用于将数据写入 TensorBoard 可读取的 `.event` 文件中。

| 参数                        | 含义                           |
|---------------------------|------------------------------|
| <code>log_dir</code>      | 保存路径                         |
| <code>flush_secs=1</code> | 每1秒就刷新一次日志（即时可见）             |
| <code>max_queue=1</code>  | 不使用队列缓存，数据写入尽可能立即进行，防止数据延迟出现 |

- ```

1     seed = params['seed']
2     np.random.seed(seed)
3     torch.manual_seed(seed)

```

设置 NumPy 的随机种子，使得：

```
1 np.random.rand(), np.random.randint(), np.random.permutation()
```

等 NumPy 的随机函数每次运行时的结果**固定**，不会随着时间或运行顺序变化。

设置 **PyTorch** 的随机种子，使得：

```
1 pythonCopyEdittorch.rand(), torch.randn(), torch.randint()
2 torch.nn.init.xavier_uniform_()
```

这些涉及到模型参数初始化或训练过程中的随机操作时，每次运行都保持一致。

- 使用 OpenAI Gym 创建一个环境并设定随机种子

`render_mode` 是 Gym API 从 v0.26+ 引入的新参数，用于指定渲染方式：

- `'human'`: 实时窗口可视化
- `'rgb_array'`: 以 NumPy 格式返回帧图像（适用于视频记录）
- `None`: 不渲染（节省资源）

```
1 env = gym.make(params['env_name'], render_mode=None)
2 env.reset(seed=seed)
```

- ```
1 assert isinstance(env.action_space, gym.spaces.Box), "Environment must be
   continuous"#连续动作空间
2 # Observation and action sizes
3 ob_dim = env.observation_space.shape[0]
4 ac_dim = env.action_space.shape[0]
```

断言（assert）用于确保当前环境的动作空间是 `gym.spaces.Box` 类型。

`gym.spaces.Box` 表示的是 **连续动作空间**，比如动作是实数，如 `[1.3, -0.2]`。

如果不是（比如是离散动作空间 `gym.spaces.Discrete`），就会报错。

- ```
1 # simulation timestep, will be used for video saving
2 if 'model' in dir(env):
3     fps = 1/env.model.opt.timestep
4 else:
5     fps = env.env.metadata['render_fps']
```

**获取环境的视频帧率（fps, frames per second）**，用于后续在 TensorBoard 中保存视频。

判断 `env`（环境）对象是否包含 `model` 这个属性。

通常 `model` 是 **Mujoco** 环境特有的属性。

如果环境是基于 MuJoCo（如 `HalfCheetah-v4`, `Ant-v4`），那么它们内部有一个低层次的 `model`，包含仿真的物理属性。

对于非 Mujoco 环境（如 Atari、CartPole 等），一般在 `env.metadata` 字典里直接定义了渲染帧率。

`render_fps` 就是推荐的可视化帧率

- `replay_buffer = ReplayBuffer(params['max_replay_buffer_size'])`

`ReplayBuffer` 类是一个用来**缓存 agent 与环境交互的样本**的模块，结构一般是一个大表格



用于打破样本之间的时间相关性，提升训练稳定性（如 DQN、SAC）。

```
1 def add_rollouts(self, paths, concat_rew=True):
2
3     # add new rollouts into our list of rollouts
4     for path in paths:
5         self.paths.append(path)
6
7     # convert new rollouts into their component arrays, and append them onto
8     # our arrays
9     observations, actions, rewards, next_observations, terminals = (
10         convert_listofrollouts(paths, concat_rew))
11
12     if self.obs is None:
13         self.obs = observations[-self.max_size:]
14         self.acs = actions[-self.max_size:]
15         self.rews = rewards[-self.max_size:]
16         self.next_obs = next_observations[-self.max_size:]
17         self.terminals = terminals[-self.max_size:]
18     else:
19         self.obs = np.concatenate([self.obs, observations])[-self.max_size:]
20         self.acs = np.concatenate([self.acs, actions])[-self.max_size:]
21         if concat_rew:
22             self.rews = np.concatenate(
23                 [self.rews, rewards]
24             )[-self.max_size:]
25         else:
26             if isinstance(rewards, list):
27                 self.rews += rewards
28             else:
29                 self.rews.append(rewards)
30             self.rews = self.rews[-self.max_size:]
31         self.next_obs = np.concatenate(
32             [self.next_obs, next_observations]
33         )[-self.max_size:]
34         self.terminals = np.concatenate(
35             [self.terminals, terminals]
36         )[-self.max_size:]
```

`concat_rew=True`：是否将 reward 拼接成一维数组（方便 supervised learning）

行为	说明
<code>self.rews += rewards</code>	批量扩展
<code>self.rews.append(rewards)</code>	单个追加
<code>self.rews = self.rews[-self.max_size:]</code>	剪裁到最大容量

调用 `convert_listofrollouts()` 将多个 path 中的 `obs`、`act` 等字段分别拼接成统一的大数组

```
1 def convert_listofrollouts(paths, concat_rew=True):
```

```

2     """
3     Take a list of rollout dictionaries
4     and return separate arrays,
5     where each array is a concatenation of that array from across the rollouts
6     """
7     observations = np.concatenate([path["observation"] for path in paths])
8     actions = np.concatenate([path["action"] for path in paths])
9     if concat_rew:
10         rewards = np.concatenate([path["reward"] for path in paths])
11     else:
12         rewards = [path["reward"] for path in paths]
13     next_observations = np.concatenate([path["next_observation"] for path in
paths])
14     terminals = np.concatenate([path["terminal"] for path in paths])
15     return observations, actions, rewards, next_observations, terminals

```

参数 <code>concat_rew</code>	输出结构	使用场景说明
<code>True</code>	<code>np.array (N,)</code>	模型训练（统一结构）
<code>False</code>	<code>List[List]</code>	日志、评估（保持轨迹结构）

- `LoadedGaussianPolicy` 是一个类，表示 一个从文件加载的高斯策略模型，通常是预训练好的专家模型。

```

1     # decide if videos should be rendered/logged at this iteration
2     log_video = ((itr % params['video_log_freq'] == 0) and (params['video_log_freq']
!= -1))
3     # decide if metrics should be logged
4     log_metrics = (itr % params['scalar_log_freq'] == 0)

```

设置`video_log_freq`和`scalar_log_freq`判断是否记录日志

```

1     if itr == 0:
2         # BC training from expert data.
3         paths = pickle.load(open(params['expert_data'], 'rb'))
4         envsteps_this_batch = 0
5     else:
6         # DAGGER training from sampled data relabeled by expert
7         assert params['do_dagger']
8         # TODO: collect `params['batch_size']` transitions
9         # HINT: use utils.sample_trajectories
10        # TODO: implement missing parts of utils.sample_trajectory
11        paths, envsteps_this_batch = utils.sample_trajectory(
12            env, actor, params['batch_size'], params['ep_len'], render=log_video)
13
14        # relabel the collected obs with actions from a provided expert policy
15        if params['do_dagger']:
16            print("\nRelabelling collected observations with labels from an expert
policy...")
17
18        # TODO: relabel collected observations (from our policy) with labels
from expert policy

```

```

19         # HINT: query the policy (using the get_action function) with paths[i]
    ["observation"]
20         # and replace paths[i]["action"] with these expert labels
21         for i in range(len(paths)):
22             paths[i]["action"] = expert_policy.get_action(
23                 paths[i]["observation"])
24             #利用专家数据替代

```

DAGGER的**核心**：利用当前策略收集状态，但用**专家指导动作**，保证训练数据的质量。

调用sample\_trajectories去重新采样包括obs和action，然后利用expert\_policy.get\_action去重新得到专家指导的action

- ```

1         for _ in range(params['num_agent_train_steps_per_iter']):
2
3             # TODO: sample some data from replay_buffer
4             # HINT1: how much data = params['train_batch_size']
5             # HINT2: use np.random.permutation to sample random indices
6             # HINT3: return corresponding data points from each array (i.e., not
different indices from each array)
7             # for imitation learning, we only need observations and actions.
8             ob_batch, ac_batch = replay_buffer.sample(
9                 params['train_batch_size'])
10
11             # use the sampled data to train an agent
12             train_log = actor.update(ob_batch, ac_batch)
13             training_logs.append(train_log)

```

调用sample函数随机选取数据进行训练，train\_log 是一个字典，可能包含update返回的 loss, accuracy 等训练信息，用于后续记录

```

1     def sample(self, batch_size):
2         """Sample given batch size of observations and actions. """
3         indices = np.random.randint(0, len(self.acs), size=(batch_size,))
4         return self.obs[indices], self.acs[indices]

```

从 [0, len(self.acs)) 范围内随机选择 batch\_size 个整数索引。

- ```

1         if log_video:
2             # save eval rollouts as videos in tensorboard event file
3             print('\nCollecting video rollouts eval')
4             eval_video_paths = utils.sample_n_trajectories(
5                 env, actor, MAX_NVIDEO, MAX_VIDEO_LEN, True)
6
7             # save videos
8             if eval_video_paths is not None:
9                 logger.log_paths_as_videos(
10                     eval_video_paths, itr,
11                     fps=fps,
12                     max_videos_to_save=MAX_NVIDEO,
13                     video_title='eval_rollouts')
14

```

重新采样几条 `trajectory`，并将 `render=True` 打开，采集 `image_obs` 并写入日志文件，观察actor是否学会

```
1 def log_paths_as_videos(self, paths, step, max_videos_to_save=2, fps=10,
2   video_title='video'): #从路径中记录视频
3
4   # reshape the rollouts
5   videos = [np.transpose(p['image_obs'], [0, 3, 1, 2]) for p in paths]
6
7   # max rollout length
8   max_videos_to_save = np.min([max_videos_to_save, len(videos)])
9   max_length = videos[0].shape[0]
10  for i in range(max_videos_to_save):
11      if videos[i].shape[0]>max_length:
12          max_length = videos[i].shape[0] #选取最长的视频
13
14  # pad rollouts to all be same length
15  for i in range(max_videos_to_save):
16      if videos[i].shape[0]<max_length:
17          padding = np.tile([videos[i][-1]], (max_length-
18  videos[i].shape[0],1,1,1))
19          videos[i] = np.concatenate([videos[i], padding], 0)
20
21  # log videos to tensorboard event file
22  videos = np.stack(videos[:max_videos_to_save], 0)
23  self.log_video(videos, video_title, step, fps=fps)
```

```
videos = [np.transpose(p['image_obs'], [0, 3, 1, 2]) for p in paths]:
```

格式	说明	用于
<code>[T, H, W, C]</code>	原始格式，常见于 OpenCV/Numpy	环境返回的图像序列
<code>[T, C, H, W]</code>	PyTorch/TF 格式，适用于神经网络	<b>TensorBoardX 日志</b> 、网络输入

```
padding = np.tile([videos[i][-1]], (max_length-videos[i].shape[0],1,1,1)):
```

- `videos[i][-1]`：当前视频最后一帧（shape: `[C, H, W]`）。
- `[videos[i][-1]]`：变成 `[1, C, H, W]`，便于 `tile` 复制。
- `np.tile(..., (N, 1, 1, 1))`：沿第0维（时间轴）复制最后一帧，使得视频长度从当前长度补齐到 `max_length`。

`videos[i] = np.concatenate([videos[i], padding], 0)` 从第0个维度进行拼接，即通过时间帧数

`videos = np.stack(videos[:max_videos_to_save], 0):np.stack(..., 0)` 的作用是把些视频沿着一个新的第0维度堆叠起来，得到一个形状为 `[N, T, C, H, W]` 的五维数组

- 1 `if log_metrics:`  
2 `# save eval metrics`  
3 `print("\nCollecting data for eval...")`  
4 `eval_paths, eval_envsteps_this_batch = utils.sample_trajectories(`  
5 `env, actor, params['eval_batch_size'], params['ep_len'])`  
6

```

7     logs = utils.compute_metrics(paths, eval_paths)
8     # compute additional metrics
9     logs.update(training_logs[-1]) # Only use the last log for now
10    logs["Train_EnvStepsSoFar"] = total_envsteps
11    logs["TimeSinceStart"] = time.time() - start_time
12    if itr == 0:
13        logs["Initial_DataCollection_AverageReturn"] =
logs["Train_AverageReturn"]
14
15    # perform the logging
16    for key, value in logs.items():
17        print('{} : {}'.format(key, value))
18        logger.log_scalar(value, key, itr)
19    print('Done logging...\n\n')
20
21    logger.flush()
22
23    if params['save_params']:
24        print('\nsaving agent params')
25        actor.save('{} /policy_itr_{}.pt'.format(params['logdir'], itr))
26

```

`logs.update(training_logs[-1]) # Only use the last log for now` 只取最后一次训练日志是为了简洁、准确地反映当前训练状态，方便后续分析和展示

`logs["Initial_DataCollection_AverageReturn"] = logs["Train_AverageReturn"]` 记录训练开始时（第0次迭代）从专家数据收集到的初始平均回报

```

1     def log_scalar(self, scalar, name, step_):
2         self._summ_writer.add_scalar('{} {}'.format(name), scalar, step_

```

利用写好的log\_scalar把数据写入tensorboard

`logger.flush()` 确保所有日志数据都写入硬盘

- ```

1     import argparse
2     parser = argparse.ArgumentParser()
3     parser.add_argument('--env_name', '-env', type=str, help=f'choices: {"",
".join(MJ_ENV_NAMES)}', required=True)
4     args = parser.parse_args()
5
6     # convert args to dictionary
7     params = vars(args)

```

导入 Python 标准库中的 `argparse` 模块。它用于从命令行解析参数，让你可以很方便地定义程序运行时的参数选项

`help` 是帮助信息，告诉用户可选项是 `MJ_ENV_NAMES` 列表里的字符串，用逗号连接显示

`required=True` 表示这个参数是必须要传的，否则程序会报错

`args = parser.parse_args()`

这行代码会从命令行解析参数，返回一个 `Namespace` 对象，里面的每个属性对应一个命令行参数及其值

`vars()` 是 Python 内置函数，用于将对象转换为字典。

它会吧 `args` 中的属性和值，转换成一个标准的 Python 字典

```
1 data_path = os.path.join(os.path.dirname(os.path.realpath(__file__)),  
2 ' ../ ../data')  
3 if not (os.path.exists(data_path)):  
4     os.makedirs(data_path)  
5     logdir = logdir_prefix + args.exp_name + '_' + args.env_name + '_' +  
6     time.strftime("%d-%m-%Y_%H-%M-%S")  
7     logdir = os.path.join(data_path, logdir)  
8     params['logdir'] = logdir  
9     if not (os.path.exists(logdir)):  
10         os.makedirs(logdir)
```

可以通过`args.name`调用，`time.strftime("%d-%m-%Y_%H-%M-%S")`：获取当前时间，并格式化为字符串，格式为“日-月-年\_时-分-秒”

## switchdagger

在一个具有时间跨度  $T$  的离散马尔可夫决策过程 (MDP) 中，存在一个专家策略  $\pi^*$ 。在算法的每个迭代步骤  $n=1, \dots, N$  中，都有一个当前策略  $\pi^n$ ；当基于该策略生成轨迹时，会在某个随机时间步将控制权转交给专家，并由专家完成轨迹的剩余部分，在  $\pi^n$  移交给专家，我们称为  $\tilde{\pi}^n$

$\hat{\pi}^n \leftarrow \text{fit to expert actions } \pi^*(s) \text{ across } s \sim p_{\tilde{\pi}^{n-1}}$

$\tilde{\pi}^n \leftarrow S^{X_n}(\hat{\pi}^n, \tilde{\pi}^{n-1})$  where  $X_n + 1 \sim \text{Geom}(1 - \alpha)$

1. 在  $n-1$  的分布下拟合专家策略
2.  $n+1$  步及以后每一步是否移交专家控制符合几何分布  $\text{Geom}(1 - \alpha)$ ，参数是移交概率

### 第一问

设定：其中  $t$  表示一共  $t$  时间跨度（从  $\tilde{\pi}^n$  之后）， $n$  表示考虑第  $n$  步（初始）是否移交

$$A(0, n) = 0,$$

$$A(t, 0) = 0,$$

$$A(t, n) = \alpha \epsilon t + \alpha(1 - \epsilon)A(t - 1, n) + (1 - \alpha)A(t, n - 1)$$

第三个式子：第一项不移交时犯错的cost（一步错步步错所以\* $t$ ），第二项不移交但这一步不犯错故时间跨步缩小，第三项表示第  $n$  步已移交，所以回退一个  $n-1$ （需要从起始开始考虑）

### 第二问

归纳法：

1. 当  $n=0$  时成立
2. 假设  $n=k$  成立
3.  $C(\tilde{\pi}^{n+1}) \leq \epsilon \alpha T + (1 - \alpha)C(\tilde{\pi}^n)$

右式第一项表示不移交专家，后一项表示移交专家回退一步

故得证

### 第三问

$\pi^n$ 与 $\tilde{\pi}^n$ 的成本差异源于 $X^*$ 的取值：

- 若 $X^* > T$ ：由于 MDP 的时间跨度为 $T$ ，无额外成本。
- 若 $X^* \leq T$ ：会产生额外的犯错期望,最大为1

$$T \cdot Pr[X^* \leq T] \leq T \cdot e^{\frac{-n}{(1-\alpha)T}}$$

$$\text{故有 } C(\pi^n) \leq C(\tilde{\pi}^n) + T e^{\frac{-n}{(1-\alpha)T}}$$

### 第四问

$$C(\pi^n) \leq T n \alpha \varepsilon + T e^{\frac{-n}{(1-\alpha)T}}$$

$$\text{取 } \alpha = \frac{1}{\log(1/\varepsilon)}, \quad N = O(T \log(1/\varepsilon))$$

结合 $\alpha \leq 1/T$ 即可